

TSF Final Assignment

Murat Tirkeshov

Kaggle competition

```
In [62]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_pacf, plot_acf
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.arima.model import ARIMA
from pmdarima import auto_arima
from sklearn.metrics import mean_absolute_error, mean_squared_error
from statsmodels.tsa.arima_process import ArmaProcess
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder, StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.stats.diagnostic import acorr_ljungbox, acorr_breus
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.api import VAR
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.stattools import arma_order_select_ic
from arch import arch_model
from statsmodels.tsa.stattools import kpss
from statsmodels.tsa.ar_model import AutoReg
from statsmodels.tsa.deterministic import CalendarFourier, DeterministicProcess
from tbats import TBATS
from statsmodels.stats.diagnostic import acorr_ljungbox
from scipy.stats import norm
import pmdarima as pm
from prophet import Prophet
from itertools import product
from statsmodels.tsa.stattools import grangercausalitytests
import itertools
import plotly.express as px

from prophet import Prophet

import tensorflow as tf
from tensorflow.keras.layers import Dense, Input, GlobalMaxPooling1D
from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Model
from sklearn.metrics import mean_absolute_percentage_error
from tensorflow.keras.callbacks import ModelCheckpoint
from sklearn.metrics import mean_absolute_percentage_error
import seaborn as sns
```

```
from statsmodels.stats.diagnostic import acorr_ljungbox
import scipy.stats as stats
import statsmodels.api as sm
import warnings
warnings.filterwarnings('ignore')
```

Data

Merging given data sets

In [3]:

```
df = pd.read_csv('store-sales-time-series-forecasting/train.csv', p
submission = pd.read_csv('store-sales-time-series-forecasting/test.
stores = pd.read_csv('store-sales-time-series-forecasting/stores.cs
oil = pd.read_csv('store-sales-time-series-forecasting/oil.csv', pa
transactions = pd.read_csv('store-sales-time-series-forecasting/tra
holidays = pd.read_csv('store-sales-time-series-forecasting/holiday
```

In [4]:

```
df.head()
```

Out[4]:

	id	date	store_nbr	family	sales	onpromotion
0	0	2013-01-01	1	AUTOMOTIVE	0.0	0
1	1	2013-01-01	1	BABY CARE	0.0	0
2	2	2013-01-01	1	BEAUTY	0.0	0
3	3	2013-01-01	1	BEVERAGES	0.0	0
4	4	2013-01-01	1	BOOKS	0.0	0

In [5]:

```
df = df.merge(stores, on='store_nbr', how='left')
submission = submission.merge(stores, on='store_nbr', how='left')
df.head()
```

Out[5]:

	id	date	store_nbr	family	sales	onpromotion	city	state
0	0	2013-01-01	1	AUTOMOTIVE	0.0	0	Quito	Pichincha
1	1	2013-01-01	1	BABY CARE	0.0	0	Quito	Pichincha
2	2	2013-01-01	1	BEAUTY	0.0	0	Quito	Pichincha
3	3	2013-01-01	1	BEVERAGES	0.0	0	Quito	Pichincha
4	4	2013-01-01	1	BOOKS	0.0	0	Quito	Pichincha

In [6]:

```
df = df.merge(transactions, on=['date', 'store_nbr'], how='left')
submission = submission.merge(transactions, on=['date', 'store_nbr']
df.head()
```

	id	date	store_nbr	family	sales	onpromotion	city	state
0	0	2013-01-01	1	AUTOMOTIVE	0.0	0	Quito	Pichincha
1	1	2013-01-01	1	BABY CARE	0.0	0	Quito	Pichincha
2	2	2013-01-01	1	BEAUTY	0.0	0	Quito	Pichincha
3	3	2013-01-01	1	BEVERAGES	0.0	0	Quito	Pichincha
4	4	2013-01-01	1	BOOKS	0.0	0	Quito	Pichincha

```
In [7]: oil = oil.set_index('date').resample('D').mean().interpolate()
oil.reset_index(inplace=True)
df = df.merge(oil, on='date', how='left')
submission = submission.merge(oil, on='date', how='left')
df.head()
```

	id	date	store_nbr	family	sales	onpromotion	city	state
0	0	2013-01-01	1	AUTOMOTIVE	0.0	0	Quito	Pichincha
1	1	2013-01-01	1	BABY CARE	0.0	0	Quito	Pichincha
2	2	2013-01-01	1	BEAUTY	0.0	0	Quito	Pichincha
3	3	2013-01-01	1	BEVERAGES	0.0	0	Quito	Pichincha
4	4	2013-01-01	1	BOOKS	0.0	0	Quito	Pichincha

```
In [8]: def process_holidays(df):
    df = df.copy()
    df['is_holiday'] = 1
    df = df[['date', 'is_holiday', 'type', 'locale']]
    df = df.drop_duplicates('date')
    df = df.pivot_table(index='date',
                         values='is_holiday',
                         aggfunc='max').reset_index()
    return df

holiday_features = process_holidays(holidays)
df = df.merge(holiday_features, on='date', how='left')
submission = submission.merge(holiday_features, on='date', how='lef

df.head()
```

Out[8]:

	id	date	store_nbr	family	sales	onpromotion	city	state
0	0	2013-01-01	1	AUTOMOTIVE	0.0	0	Quito	Pichincha
1	1	2013-01-01	1	BABY CARE	0.0	0	Quito	Pichincha
2	2	2013-01-01	1	BEAUTY	0.0	0	Quito	Pichincha
3	3	2013-01-01	1	BEVERAGES	0.0	0	Quito	Pichincha
4	4	2013-01-01	1	BOOKS	0.0	0	Quito	Pichincha

Missing values

In [9]:

```
df['transactions'] = df['transactions'].fillna(df['transactions'].mean())
submission['transactions'] = submission['transactions'].fillna(submission['transactions'].mean())
df['dcoilwtico'] = df['dcoilwtico'].fillna(df['dcoilwtico'].mean())
submission['dcoilwtico'] = submission['dcoilwtico'].fillna(submission['dcoilwtico'].mean())
```

Adding aditional variables

In [10]:

```
df['day'] = df['date'].dt.day
df['month'] = df['date'].dt.month
df['year'] = df['date'].dt.year
df['day_of_week'] = df['date'].dt.dayofweek
df['is_weekend'] = (df['day_of_week'] >= 5).astype(int)
df['is_weekend_and_holiday'] = ((df['is_weekend'] == 1) & (df['is_holiday'] == 1))
submission['day'] = submission['date'].dt.day
submission['month'] = submission['date'].dt.month
submission['year'] = submission['date'].dt.year
submission['day_of_week'] = submission['date'].dt.dayofweek
submission['is_weekend'] = (submission['day_of_week'] >= 5).astype(int)
submission['is_weekend_and_holiday'] = ((submission['is_weekend'] == 1) & (submission['is_holiday'] == 1))
```


In [11]:

```
df['is_holiday'] = df['is_holiday'].fillna(0)
submission['is_holiday'] = submission['is_holiday'].fillna(0)
```


Data Visualization

In [12]:

```
daily = df.groupby('date')['sales'].sum().reset_index()
plt.figure(figsize=(12,4))
plt.plot(daily['date'], daily['sales'])
plt.title('Total Daily Sales (All Stores & Families)')
plt.xlabel('Date')
plt.ylabel('Sales')
plt.tight_layout()
plt.show()
```




```
In [63]: mask = (df['store_nbr']==1) & (df['family']=='GROCERY I')
ts = df[mask].set_index('date')['sales'].asfreq('D').fillna(0)
decomp = seasonal_decompose(ts, model='additive', period=7)
fig = decomp.plot()
fig.set_size_inches(12,8)
plt.tight_layout()
plt.title('Seasonal Decomposition of Sales for Store 1 - GROCERY I')
plt.show()
```




```
In [14]: plt.figure(figsize=(12,4))
sns.histplot(df['sales'], bins=100, kde=True)
plt.xlim(0, df['sales'].quantile(0.99))
plt.title('Distribution of Daily Sales')
plt.xlabel('Sales')
plt.show()

plt.figure(figsize=(12,4))
sns.histplot(np.log1p(df['sales']), bins=100, kde=True)
plt.title('Distribution of log1p(Sales)')
plt.xlabel('log1p(Sales)')
plt.show()
```




```
In [15]: plt.figure(figsize=(8,4))
sns.boxplot(x='day_of_week', y='sales', data=df)
plt.ylim(0, df['sales'].quantile(0.95))
plt.title('Sales by Day of Week (0=Mon ... 6=Sun)')
plt.show()
```




```
In [16]: pivot = df.groupby(['year', 'month'])['sales'].mean().unstack(level=0)
plt.figure(figsize=(10,6))
sns.heatmap(pivot, annot=True, fmt=".0f", cmap='YlGnBu')
plt.title('Average Monthly Sales by Year')
plt.ylabel('Month')
plt.xlabel('Year')
```


```
plt.show()
```




```
In [17]: promo = df.copy()
promo['onpromotion'] = promo['onpromotion'] > 0
mean_sales = promo.groupby('onpromotion')['sales'].mean().reset_index()
plt.figure(figsize=(6,4))
sns.barplot(x='onpromotion', y='sales', data=mean_sales)
plt.xticks([0,1], ['No Promo', 'Promo'])
plt.title('Average Sales: Promo vs No Promo')
plt.ylabel('Mean Sales')
plt.show()
```



```
In [18]: num_cols = ['sales','onpromotion','transactions','dcoilwtico','day_
corr = df[num_cols].corr()
plt.figure(figsize=(8,6))
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```




```
In [19]: family_sales = df.groupby('family')['sales'].sum().sort_values(ascending=True)
plt.figure(figsize=(10,5))
sns.barplot(x=family_sales.values, y=family_sales.index)
plt.title('Top 10 Product Families by Total Sales')
plt.xlabel('Total Sales')
plt.ylabel('Family')
plt.show()
```




```
In [20]: cluster_sales = df.groupby('cluster')['sales'].sum().sort_values(ascending=True)
plt.figure(figsize=(8,4))
sns.barplot(x=cluster_sales.index, y=cluster_sales.values)
```

```
plt.title('Total Sales by Store Cluster')
plt.xlabel('Cluster')
plt.ylabel('Sales')
plt.show()
```



```
In [21]: a = df[["store_nbr", "sales"]]
a["ind"] = 1
a["ind"] = a.groupby("store_nbr").ind.cumsum().values
a = pd.pivot(a, index = "ind", columns = "store_nbr", values = "sales")
mask = np.triu(a.corr())
plt.figure(figsize=(20, 20))
sns.heatmap(a,
            annot=True,
            fmt='.1f',
            cmap='coolwarm',
            square=True,
            mask=mask,
            linewidths=1,
            cbar=False)
plt.title("Correlations among stores", fontsize = 20)
plt.show()
```

Correlations among stores




ACF and PACF for store 1

```
In [22]: agg1 = (
    df[(df['store_nbr'] == 1) & df['sales'].notnull()]
    .groupby(['family', 'date'])['sales']
    .mean()
    .reset_index()
)





# 3. Loop over each family in store 1
for fam, grp in agg1.groupby('family'):
    ts = grp.set_index('date')['sales']


    try:
        fig, axes = plt.subplots(1, 2, figsize=(15, 4))
        sm.graphics.tsa.plot_acf(ts, lags=40, ax=axes[0])
        axes[0].set_title(f'Store 1 - {fam}\nACF')
        sm.graphics.tsa.plot_pacf(ts, lags=40, ax=axes[1])
        axes[1].set_title(f'Store 1 - {fam}\nPACF')
        plt.tight_layout()
        plt.show()
    except:
        pass
```


```
except Exception as e:
    print(f"Could not plot family {fam}: {e}")
```





```
/opt/anaconda3/lib/python3.12/site-packages/statsmodels/regression/linear_model.py:1491: ValueWarning: Matrix is singular. Using pinv.
  warnings.warn("Matrix is singular. Using pinv.", ValueWarning)
```














Model training

```
In [23]: train = df[df['date'] <= '2017-07-31']
test = df[(df['date'] > '2017-07-31') & (df['date'] <= '2017-08-15')

print("Training period:", train['date'].min(), "to", train['date'].max())
print("Validation period:", test['date'].min(), "to", test['date'].max())
print("Train size:", len(train), "rows", "Validation size:", len(test))
train.head()
```

Training period: 2013-01-01 00:00:00 to 2017-07-31 00:00:00
 Validation period: 2017-08-01 00:00:00 to 2017-08-15 00:00:00
 Train size: 2974158 rows, Validation size: 26730 rows

Out [23]:

	id	date	store_nbr	family	sales	onpromotion	city	state
0	0	2013-01-01	1	AUTOMOTIVE	0.0	0	Quito	Pichincha
1	1	2013-01-01	1	BABY CARE	0.0	0	Quito	Pichincha
2	2	2013-01-01	1	BEAUTY	0.0	0	Quito	Pichincha
3	3	2013-01-01	1	BEVERAGES	0.0	0	Quito	Pichincha
4	4	2013-01-01	1	BOOKS	0.0	0	Quito	Pichincha

SAIMAX model

In [24]:

```

series_key = (1, 'GROCERY I')
train_series = train[(train['store_nbr']==series_key[0]) &
                     (train['family']==series_key[1])]
val_series = test[(test['store_nbr']==series_key[0]) &
                  (test['family']==series_key[1])]

y_train = train_series['sales']
exog_features = ['onpromotion', 'is_holiday', 'dcoilwtico', 'transa
exog_train = train_series[exog_features]

sarimax_model = SARIMAX(y_train, exog=exog_train, order=(1,1,1), se
                           enforce_stationarity=False, enforce_inverti
sarimax_result = sarimax_model.fit(disp=False)
print(sarimax_result.summary().tables[1])
n_val = len(val_series)
exog_val = val_series[exog_features]
sarimax_forecast = sarimax_result.forecast(steps=n_val, exog=exog_v
def rmsle(y_true, y_pred):
    return np.sqrt(np.mean(np.square(np.log1p(y_pred) - np.log1p(y_))

series_rmsle = rmsle(val_series['sales'].values, sarimax_forecast.v
print(f"SARIMAX validation RMSLE for store {series_key[0]}, family

```

```
=====
=====
          coef    std err      z     P>|z|    [0.025
0.975]
-----
onpromotion      5.3473    0.610     8.769    0.000    4.152
6.542
is_holiday     -80.8851   32.132    -2.517    0.012   -143.862
-17.908
dcoilwtico     -33.1567   5.893    -5.627    0.000   -44.706
-21.608
transactions     1.2234    0.029    42.494    0.000    1.167
1.280
ar.L1            0.3605    0.014    26.169    0.000    0.333
0.387
ma.L1           -1.0507   0.009   -116.925    0.000   -1.068
-1.033
ar.S.L7          -0.0022   0.024    -0.093    0.926   -0.049
0.044
ma.S.L7          -0.8748   0.017   -50.302    0.000   -0.909
-0.841
sigma2         1.485e+05  2766.743    53.656    0.000  1.43e+05
1.54e+05
=====
=====
```

SARIMAX validation RMSLE for store 1, family GROCERY I: 0.1419

```
In [56]: results = []
combinations = train[['store_nbr', 'family']].drop_duplicates()

for _, row in combinations.iterrows():
    store, family = row['store_nbr'], row['family']

    train_series = train[(train['store_nbr'] == store) & (train['fa
val_series = test[(test['store_nbr'] == store) & (test['famil

    if len(train_series) < 30 or len(val_series) == 0:
        continue

    try:
        y_train = train_series['sales']
        exog_train = train_series[exog_features]
        exog_val = val_series[exog_features]

        model = SARIMAX(
            y_train, exog=exog_train,
            order=(1, 1, 1), seasonal_order=(1, 1, 1, 7),
            enforce_stationarity=False, enforce_invertibility=False
        )
        result = model.fit(disp=False)
        forecast = result.forecast(steps=len(val_series), exog=exog

        rmsle_score = rmsle(val_series['sales'].values, forecast.va
        results.append({'store_nbr': store, 'family': family, 'rmsl
    except Exception as e:
```

```

    print(f"Skipped ({store}, {family}) due to error: {e}")
    continue

results_df = pd.DataFrame(results)
results_df.to_csv("sarimax_rmsle_results.csv", index=False)
print(results_df.sort_values('rmsle').head(10))

```

	store_nbr	family	rmsle
829	32	BOOKS	0.0
1354	47	BABY CARE	0.0
679	28	LADIESWEAR	0.0
37	10	BOOKS	0.0
877	33	LADIESWEAR	0.0
664	28	BOOKS	0.0
1240	43	LADIESWEAR	0.0
250	16	LADIESWEAR	0.0
1141	40	LADIESWEAR	0.0
928	35	BOOKS	0.0

In [61]:

```

results = pd.DataFrame(results)
overall_rmsle = np.sqrt(np.mean(results['rmsle'] ** 2))

print("Overall RMSLE:", overall_rmsle)

```

Overall RMSLE: 0.42663089438997875

CatBoost

In [50]:

```

X_train = train.drop(columns=['date', 'sales'])
X_test = test.drop(columns=['date', 'sales'])
y_train = train['sales']
y_test = test['sales']

from catboost import CatBoostRegressor
cat_feature_names = ['family', 'city', 'state', 'type']
cat_feature_indices = [X_train.columns.get_loc(col) for col in cat_


model = CatBoostRegressor(
    iterations=600,
    learning_rate=0.1,
    depth=7,
    verbose=100
)

model.fit(
    X_train,
    y_train,
    cat_features=cat_feature_indices
)

preds = model.predict(X_test)

```

```
0:      learn: 1054.9797107      total: 638ms      remaining: 6m 21s
100:     learn: 358.6449079      total: 47.3s      remaining: 3m 53s
200:     learn: 323.2161544      total: 1m 40s      remaining: 3m 19s
300:     learn: 301.4173992      total: 2m 37s      remaining: 2m 36s
400:     learn: 287.4052151      total: 3m 32s      remaining: 1m 45s
500:     learn: 276.0157785      total: 4m 26s      remaining: 52.7s
599:     learn: 267.6864603      total: 5m 20s      remaining: 0us
```

```
In [51]: rmsle_cat = rmsle(y_test, preds)
print(f"CatBoost validation RMSLE: {rmsle_cat:.4f}")
```

```
CatBoost validation RMSLE: 0.9997
```

```
In [ ]:
```

Deeplearning Models

ANN

```
In [43]: for col in cat_feature_names:
    le = LabelEncoder()
    X_train[col] = le.fit_transform(X_train[col].astype(str))
    X_test[col] = le.transform(X_test[col].astype(str))

num_cols = X_train.columns.difference(cat_feature_names).tolist()
scaler = StandardScaler()
X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])

input_dim = X_train.shape[1]

model_ann = Sequential([
    Dense(128, activation='relu', input_dim=input_dim),
    Dropout(0.2),
    Dense(64, activation='relu'),
    Dropout(0.1),
    Dense(1)
])
model_ann.compile(optimizer='adam', loss='mse')

history = model_ann.fit(
    X_train.values, y_train,
    validation_data=(X_test.values, y_test),
    epochs=50,
    batch_size=32,
    verbose=1
)

ann_preds = model_ann.predict(X_test.values).flatten().clip(0)
print("ANN RMSLE:", rmsle(y_test, ann_preds))
```

```
Epoch 1/50
92943/92943 ━━━━━━━━━━ 35s 375us/step - loss: 724965.4375
- val_loss: 427405.8125
```

Epoch 2/50
92943/92943  **36s** 383us/step - loss: 395747.2812
- val_loss: 325056.9375
Epoch 3/50
92943/92943  **36s** 384us/step - loss: 318025.2812
- val_loss: 225053.8281
Epoch 4/50
92943/92943  **37s** 398us/step - loss: 270169.2812
- val_loss: 214891.3594
Epoch 5/50
92943/92943  **38s** 405us/step - loss: 256036.3281
- val_loss: 199257.5781
Epoch 6/50
92943/92943  **35s** 380us/step - loss: 251677.5469
- val_loss: 201097.6875
Epoch 7/50
92943/92943  **34s** 367us/step - loss: 226579.2031
- val_loss: 186281.5000
Epoch 8/50
92943/92943  **34s** 362us/step - loss: 224704.2344
- val_loss: 173584.9688
Epoch 9/50
92943/92943  **34s** 369us/step - loss: 226496.9844
- val_loss: 171293.4531
Epoch 10/50
92943/92943  **36s** 389us/step - loss: 210443.5625
- val_loss: 167952.7812
Epoch 11/50
92943/92943  **35s** 380us/step - loss: 218922.2031
- val_loss: 192884.3750
Epoch 12/50
92943/92943  **36s** 391us/step - loss: 207073.8906
- val_loss: 155555.0156
Epoch 13/50
92943/92943  **35s** 379us/step - loss: 217547.2031
- val_loss: 157898.2812
Epoch 14/50
92943/92943  **36s** 390us/step - loss: 200611.4375
- val_loss: 192516.2656
Epoch 15/50
92943/92943  **36s** 384us/step - loss: 213970.6406
- val_loss: 161398.6250
Epoch 16/50
92943/92943  **36s** 385us/step - loss: 196884.9375
- val_loss: 155616.1250
Epoch 17/50
92943/92943  **36s** 390us/step - loss: 198663.9219
- val_loss: 193898.9844
Epoch 18/50
92943/92943  **35s** 376us/step - loss: 192877.5781
- val_loss: 246439.1875
Epoch 19/50
92943/92943  **36s** 389us/step - loss: 190941.3750
- val_loss: 145925.8750
Epoch 20/50
92943/92943  **36s** 384us/step - loss: 205003.0000

```
- val_loss: 160612.0156
Epoch 21/50
92943/92943 36s 382us/step - loss: 196211.1250
- val_loss: 166851.0938
Epoch 22/50
92943/92943 37s 392us/step - loss: 197947.5000
- val_loss: 165686.4375
Epoch 23/50
92943/92943 34s 369us/step - loss: 189226.0156
- val_loss: 173440.6875
Epoch 24/50
92943/92943 33s 359us/step - loss: 177975.5938
- val_loss: 163216.2188
Epoch 25/50
92943/92943 35s 381us/step - loss: 186733.8594
- val_loss: 147762.7812
Epoch 26/50
92943/92943 34s 363us/step - loss: 189749.2500
- val_loss: 149810.2969
Epoch 27/50
92943/92943 34s 360us/step - loss: 178990.6406
- val_loss: 153051.8750
Epoch 28/50
92943/92943 34s 370us/step - loss: 186934.2969
- val_loss: 161705.8125
Epoch 29/50
92943/92943 36s 391us/step - loss: 177098.0000
- val_loss: 135121.9844
Epoch 30/50
92943/92943 34s 364us/step - loss: 170950.6875
- val_loss: 139069.6250
Epoch 31/50
92943/92943 34s 362us/step - loss: 172360.5781
- val_loss: 148676.0469
Epoch 32/50
92943/92943 33s 359us/step - loss: 172744.6250
- val_loss: 168409.1719
Epoch 33/50
92943/92943 35s 372us/step - loss: 174352.8906
- val_loss: 149469.8594
Epoch 34/50
92943/92943 33s 360us/step - loss: 175414.9375
- val_loss: 140922.4688
Epoch 35/50
92943/92943 34s 367us/step - loss: 172557.2500
- val_loss: 146215.8438
Epoch 36/50
92943/92943 33s 357us/step - loss: 179342.8281
- val_loss: 129323.3438
Epoch 37/50
92943/92943 33s 355us/step - loss: 168879.8281
- val_loss: 148916.5938
Epoch 38/50
92943/92943 34s 366us/step - loss: 172394.5156
- val_loss: 160249.1406
Epoch 39/50
```

```

92943/92943 ━━━━━━━━━━ 36s 382us/step - loss: 186995.5625
- val_loss: 137815.1875
Epoch 40/50
92943/92943 ━━━━━━━━━━ 35s 376us/step - loss: 170306.4219
- val_loss: 139160.8750
Epoch 41/50
92943/92943 ━━━━━━━━━━ 35s 377us/step - loss: 171504.1094
- val_loss: 148385.6094
Epoch 42/50
92943/92943 ━━━━━━━━━━ 35s 378us/step - loss: 164788.8438
- val_loss: 152800.0625
Epoch 43/50
92943/92943 ━━━━━━━━━━ 35s 377us/step - loss: 178498.2188
- val_loss: 132143.8906
Epoch 44/50
92943/92943 ━━━━━━━━━━ 35s 378us/step - loss: 170506.1875
- val_loss: 136636.1875
Epoch 45/50
92943/92943 ━━━━━━━━━━ 35s 377us/step - loss: 186354.8750
- val_loss: 135458.5781
Epoch 46/50
92943/92943 ━━━━━━━━━━ 35s 377us/step - loss: 166900.4531
- val_loss: 135461.5625
Epoch 47/50
92943/92943 ━━━━━━━━━━ 35s 376us/step - loss: 162601.3438
- val_loss: 130549.3125
Epoch 48/50
92943/92943 ━━━━━━━━━━ 35s 374us/step - loss: 169067.8906
- val_loss: 131386.4844
Epoch 49/50
92943/92943 ━━━━━━━━━━ 35s 379us/step - loss: 166145.7188
- val_loss: 152522.2656
Epoch 50/50
92943/92943 ━━━━━━━━━━ 35s 376us/step - loss: 168264.9844
- val_loss: 156142.3750
836/836 ━━━━━━━━ 0s 219us/step
ANN RMSLE: 1.7919142241944601

```

CNN

```

In [ ]: full = pd.concat([train, test]) \
          .sort_values(['store_nbr', 'family', 'date']) \
          .reset_index(drop=True)

le_store = LabelEncoder()
full['store_enc'] = le_store.fit_transform(full['store_nbr'])
le_fam = LabelEncoder()
full['fam_enc'] = le_fam.fit_transform(full['family'])

seq_feats = [
    'sales', 'onpromotion', 'is_holiday',
    'dcoilwtico', 'transactions',
    'store_enc', 'fam_enc'
]

```

```
[1]:\n\nN = 14\ntrain_max_date = train['date'].max()\n\nX_tr_seq, y_tr_seq = [], []\nX_te_seq, y_te_seq = [], []\n\nfor _, _, grp in full.groupby(['store_nbr', 'family']):\n    grp = grp.sort_values('date').reset_index(drop=True)\n    for i in range(N, len(grp)):\n        window = grp.loc[i-N:i-1, seq_feats].values\n        target_date = grp.loc[i, 'date']\n        target_sale = grp.loc[i, 'sales']\n\n        if target_date <= train_max_date:\n            X_tr_seq.append(window); y_tr_seq.append(target_sale)\n        else:\n            X_te_seq.append(window); y_te_seq.append(target_sale)\n\nX_tr_seq = np.array(X_tr_seq)\ny_tr_seq = np.array(y_tr_seq)\nX_te_seq = np.array(X_te_seq)\ny_te_seq = np.array(y_te_seq)\n\nprint("train seq:", X_tr_seq.shape, "test seq:", X_te_seq.shape)
```

train seq: (2949210, 14, 7) test seq: (26730, 14, 7)

```
-----
NameError                                 Traceback (most recent call
l last)
Cell In[57], line 53
    46 print("train seq:", X_tr_seq.shape, "test seq:", X_te_seq.sh
ape)
    48 #
-
-
-
49 # 2) Define & train the CNN
50 model_cnn = Sequential([
51     Conv1D(32, 3, activation='relu', input_shape=(N, len(seq
_feats))),
52     MaxPooling1D(2),
53     Flatten(),
54     Dense(64, activation='relu'),
55     Dense(1)  # linear output
56 ])
57 model_cnn.compile(optimizer='adam', loss='mse')
59 history = model_cnn.fit(
60     X_tr_seq, y_tr_seq,
61     validation_data=(X_te_seq, y_te_seq),
62     ...
63     verbose=1
64 )
65 )

NameError: name 'Flatten' is not defined
```


```
In [59]: model_cnn = Sequential([
    Conv1D(32, 3, activation='relu', input_shape=(N, len(seq_feats)),
    MaxPooling1D(2),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(1)
])
model_cnn.compile(optimizer='adam', loss='mse')

history = model_cnn.fit(
    X_tr_seq, y_tr_seq,
    validation_data=(X_te_seq, y_te_seq),
    epochs=20,
    batch_size=64,
    verbose=1
)

cnn_preds = model_cnn.predict(X_te_seq).flatten().clip(0)
def rmsle(y_true, y_pred):
    return np.sqrt(np.mean((np.log1p(y_pred)-np.log1p(y_true))**2))

print("CNN RMSLE:", rmsle(y_te_seq, cnn_preds))
```

Epoch 1/20
46082/46082 ━━━━━━━━ 32s 689us/step - loss: 146681.0156

```
- val_loss: 90892.6797
Epoch 2/20
46082/46082  29s 625us/step - loss: 76827.0547 -
val_loss: 64996.5391
Epoch 3/20
46082/46082  28s 606us/step - loss: 82135.6719 -
val_loss: 57614.9102
Epoch 4/20
46082/46082  28s 606us/step - loss: 72212.2578 -
val_loss: 62621.5430
Epoch 5/20
46082/46082  28s 607us/step - loss: 77133.5234 -
val_loss: 59538.0586
Epoch 6/20
46082/46082  28s 612us/step - loss: 70197.7500 -
val_loss: 72716.6484
Epoch 7/20
46082/46082  1409s 31ms/step - loss: 73327.0391
- val_loss: 53333.1328
Epoch 8/20
46082/46082  27s 575us/step - loss: 70030.2266 -
val_loss: 79379.4531
Epoch 9/20
46082/46082  27s 593us/step - loss: 76086.3125 -
val_loss: 55868.1289
Epoch 10/20
46082/46082  28s 617us/step - loss: 67867.5391 -
val_loss: 50817.9297
Epoch 11/20
46082/46082  28s 616us/step - loss: 59951.6719 -
val_loss: 57692.3242
Epoch 12/20
46082/46082  28s 613us/step - loss: 68322.5938 -
val_loss: 59911.7773
Epoch 13/20
46082/46082  28s 616us/step - loss: 73027.4141 -
val_loss: 54790.2188
Epoch 14/20
46082/46082  28s 615us/step - loss: 67984.0312 -
val_loss: 56354.3984
Epoch 15/20
46082/46082  29s 619us/step - loss: 77199.9453 -
val_loss: 57063.7305
Epoch 16/20
46082/46082  29s 620us/step - loss: 66860.2734 -
val_loss: 57783.9414
Epoch 17/20
46082/46082  28s 616us/step - loss: 63030.0547 -
val_loss: 56720.0547
Epoch 18/20
46082/46082  30s 641us/step - loss: 56338.0195 -
val_loss: 63579.8828
Epoch 19/20
46082/46082  47s 1ms/step - loss: 65576.3906 - v
al_loss: 56965.6836
Epoch 20/20
```

```
46082/46082 ━━━━━━━━ 29s 620us/step - loss: 71273.6797 -
val_loss: 57971.0625
836/836 ━━━━━━━━ 0s 259us/step
CNN RMSLE: 1.018394351235643
```

The best model CatBoost

training submission data

```
In [ ]: features = df.drop(columns=['date', 'sales'])
cat_feature_indices = [features.columns.get_loc(col) for col in cat]
model = CatBoostRegressor(
    iterations=600,
    learning_rate=0.1,
    depth=7,
    verbose=100
)
model.fit(
    features,
    df['sales'],
    cat_features=cat_feature_indices
)
```

	learn:	total:	remaining:
0:	1056.2901431	601ms	5m 59s
100:	367.3552426	47s	3m 52s
200:	326.4408369	1m 35s	3m 10s
300:	304.3553557	2m 24s	2m 23s
400:	290.4660154	3m 12s	1m 35s
500:	279.5679376	4m 6s	48.6s
599:	270.6415999	5m 1s	0us

```
-----
TypeError                                         Traceback (most recent call
l last)
File _catboost.pyx:2547, in _catboost.get_float_feature()

File _catboost.pyx:1226, in _catboost._FloatOrNan()

File _catboost.pyx:1021, in _catboost._FloatOrNanFromString()

TypeError: Cannot convert 'b'AUTOMOTIVE'' to float
```

During handling of the above exception, another exception occurred:

```
CatBoostError                                         Traceback (most recent call
l last)
Cell In[54], line 14
    3 model = CatBoostRegressor(
    4     iterations=600,
    5     learning_rate=0.1,
    6     depth=7,
    7     verbose=100
    8 )
    9 model.fit()
```

```
10     features,
11     df['sales'],
12     cat_features=cat_feature_indices
13 )
--> 14 preds = model.predict(submission.drop(columns=['date', 'id']))

File /opt/anaconda3/lib/python3.12/site-packages/catboost/core.py:59
24, in CatBoostRegressor.predict(self, data, prediction_type, ntree_start, ntree_end, thread_count, verbose, task_type)
5922 if prediction_type is None:
5923     prediction_type = self._get_default_prediction_type()
-> 5924 return self._predict(data, prediction_type, ntree_start, ntree_end, thread_count, verbose, 'predict', task_type)

File /opt/anaconda3/lib/python3.12/site-packages/catboost/core.py:26
20, in CatBoost._predict(self, data, prediction_type, ntree_start, ntree_end, thread_count, verbose, parent_method_name, task_type)
2618 if verbose is None:
2619     verbose = False
-> 2620 data, data_is_single_object = self._process_predict_input_data(data, parent_method_name, thread_count)
2621 self._validate_prediction_type(prediction_type)
2623 predictions = self._base_predict(data, prediction_type, ntree_start, ntree_end, thread_count, verbose, task_type)

File /opt/anaconda3/lib/python3.12/site-packages/catboost/core.py:26
00, in CatBoost._process_predict_input_data(self, data, parent_method_name, thread_count, label)
2598 is_single_object = _is_data_single_object(data)
2599 if not isinstance(data, Pool):
-> 2600     data = Pool(
2601         data=[data] if is_single_object else data,
2602         label=label,
2603         cat_features=self._get_cat_feature_indices() if not
isinstance(data, FeaturesData) else None,
2604         text_features=self._get_text_feature_indices() if not
isinstance(data, FeaturesData) else None,
2605         embedding_features=self._get_embedding_feature_indices() if not
isinstance(data, FeaturesData) else None,
2606         thread_count=thread_count
2607     )
2608 return data, is_single_object

File /opt/anaconda3/lib/python3.12/site-packages/catboost/core.py:85
5, in Pool.__init__(self, data, label, cat_features, text_features,
embedding_features, embedding_features_data, column_description, pairs,
graph, delimiter, has_header, ignore_csv_quoting, weight, group_id,
group_weight, subgroup_id, pairs_weight, baseline, timestamp,
feature_names, feature_tags, thread_count, log_cout, log_cerr, data_can_be_none)
849     if isinstance(feature_names, PATH_TYPES):
850         raise CatBoostError(
851             "feature_names must be None or have non-string type when the pool is created from "
852             "python objects."
```

```

853
--> 855         self._init(data, label, cat_features, text_features,
embedding_features, embedding_features_data, pairs, graph, weight,
856                     group_id, group_weight, subgroup_id, pair
s_weight, baseline, timestamp, feature_names, feature_tags, thread_c
ount)
857 elif not data_can_be_none:
858     raise CatBoostError("'data' parameter can't be None")

File /opt/anaconda3/lib/python3.12/site-packages/catboost/core.py:14
91, in Pool._init(self, data, label, cat_features, text_features, em
bedding_features, embedding_features_data, pairs, graph, weight, gro
up_id, group_weight, subgroup_id, pairs_weight, baseline, timestamp,
feature_names, feature_tags, thread_count)
1489 if feature_tags is not None:
1490     feature_tags = self._check_transform_tags(feature_tags,
feature_names)
-> 1491 self._init_pool(data, label, cat_features, text_features, em
bedding_features, embedding_features_data, pairs, graph, weight,
1492                     group_id, group_weight, subgroup_id, pairs_w
eight, baseline, timestamp, feature_names, feature_tags, thread_coun
t)

File _catboost.pyx:4339, in _catboost._PoolBase._init_pool()

File _catboost.pyx:4391, in _catboost._PoolBase._init_pool()

File _catboost.pyx:4200, in _catboost._PoolBase._init_features_order
_layout_pool()

File _catboost.pyx:3127, in _catboost._set_features_order_data_pd_da
ta_frame()

File _catboost.pyx:2591, in _catboost.create_num_factor_data()

File _catboost.pyx:2549, in _catboost.get_float_feature()

CatBoostError: Bad value for num_feature[non_default_doc_idx=0,featu
re_idx=1]="AUTOMOTIVE": Cannot convert 'b'AUTOMOTIVE'' to float

```

```
In [ ]: X_submission = submission[features.columns]
preds = model.predict(X_submission)
sub = submission[['id']].copy()
sub['sales'] = preds.clip(0)
sub.to_csv('submission_catboost.csv', index=False)
```

Random Forest

```
In [66]: X_train_enc = train.drop(columns=['date', 'sales']).copy()
y_train = train['sales']

X_test_enc = test.drop(columns=['date', 'sales']).copy()
y_test = test['sales']
```

```

# Identify object columns (categoricals)
cat_cols = X_train_enc.select_dtypes(include=['object', 'category'])

# Label encode them
for col in cat_cols:
    le = LabelEncoder()
    X_train_enc[col] = le.fit_transform(X_train_enc[col].astype(str))
    X_test_enc[col] = le.transform(X_test_enc[col].astype(str))

# Now you can fit the RandomForest
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(random_state=42)
rf.fit(X_train_enc, y_train)

# Predict
rf_preds = rf.predict(X_test_enc)

# Evaluate
def rmsle(y_true, y_pred):
    return np.sqrt(np.mean((np.log1p(y_pred) - np.log1p(y_true))**2))

print("Random Forest RMSLE:", rmsle(y_test, rf_preds))

```

Random Forest RMSLE: 0.49665852946033556

In [67]:

```

target = df['sales']

cat_cols = features.select_dtypes(include=['object', 'category']).columns
submission_enc = X_submission.copy()

features_enc = features.copy()
for col in cat_cols:
    le = LabelEncoder()
    features_enc[col] = le.fit_transform(features_enc[col].astype(str))
    submission_enc[col] = le.transform(submission_enc[col].astype(str))

rf = RandomForestRegressor(random_state=42, n_estimators=100)
rf.fit(features_enc, target)
submission_preds = rf.predict(submission_enc)

```

In [69]:

```

submission_preds = pd.DataFrame({
    'id': submission['id'],
    'sales': submission_preds.clip(0)
})
submission_preds.to_csv('submission_pred.csv', index=False)

```

At the end i tried Random forest without parameter tuning, because i thought we were not using everything possible with

In []: